

Pranshu Gupta

<https://pranshug.com> | [linkedin.com/in/pranshug258](https://www.linkedin.com/in/pranshug258)

Email: pranshug258@outlook.com

Mobile: +1 (425) 679-2402

PROFESSIONAL EXPERIENCE

- **Microsoft AI** / *Member of Technical Staff* / Redmond, WA Aug 2025 – Present
 - Reduced production GPU footprint by **1,320 H100 GPUs** by analyzing the inference pipeline and eliminating non-essential classifiers and summarizers on the GPT-5.1 instant path.
 - Led end-to-end rollout of **GPT-5.1 instant and reasoning models** to production by designing prompt templates, constrained-sampling, chain of thought handling, and safety integrations across six partner teams; owned launch-readiness evaluations with the Response Quality squad, including **30% improvement in personality evals**.
 - Designed and shipped **Adaptive Thinking**, a dynamic system that adjusts the thinking effort per message based on intent/capability classification to control reasoning token usage.
 - Architected and shipped an autonomous **code simplification agent** that scans recent commits, scopes safe refactors via durable git baselines, and manages the full PR lifecycle through automated validation.
 - Discovered and fixed a bug in the OpenAI Chat Completions API causing function-definition parameter descriptions to be duplicated, saving **6K tokens per inference call** and reducing **TTFT by 30ms**.
 - Built **Grafana dashboards** for LLM API traffic classification across MAI and M365 Copilot, instrumenting dimensions including product, client app, agent, inference intent, license tier, and traffic type; implemented production failure-rate monitor with alerting for Bing Search API dependency.
- **Microsoft Azure AI** / *Senior Software Engineer* / Redmond, WA Jan 2025 – Aug 2025
 - Designed a distributed **token-based rate limiter** using Tiktoken for precise token counting, replacing a length-based estimator with **219% average error**; implemented pub-sub quota state synchronization across instances.
 - Contributed to an intelligent **model router** for per-request model selection based on cost/complexity signals, reducing per-inference cost by **55%**.
 - Designed **multi-region routing architecture** for the M365 LLM inference API: same-region-preferred routing with ordered fallback regions to balance GPU session affinity against compute availability spillover across global deployments.
- **Microsoft Azure Core** / *Senior Software Engineer* / Redmond, WA Mar 2024 – Jan 2025
 - Designed and shipped a token-based job throttling system with multi-bucket quota enforcement and retry-storm protection, preventing partition-level noisy-neighbor throttling in background-job infrastructure of Azure Control Plane.
 - Single-handedly designed and shipped async operation URI signing, adopted by **242 Azure service teams** and enforced for **220+ resource providers**.
- **Microsoft Azure Core** / *Software Engineer II* / Redmond, WA Jun 2020 – Mar 2024
 - Designed and shipped async operation callbacks replacing polling, reducing ARM-to-resource-provider traffic by up to **94%** (130M fewer requests/day) and improving p95 latency by **12 seconds** for Azure Container Instances.
 - Extended and maintained **Azure DevOps CI/CD pipelines** for ARM control-plane services with staged multi-region rollouts, 24-hour regional bake times, functional test gates, and rollback-via-redeploy across global Azure regions.
 - Led **Azure region teardown** end-to-end: built workflows to halt new resource creation, resolve subscription/resource-group hierarchies, and migrate live resources to healthy regions with zero data loss.
- **Microsoft Core Services** / *Software Engineer* / Hyderabad, India Jun 2017 – Aug 2019
 - Migrated customer data enrichment service to serverless design, reducing cost by **90%** and improving throughput by **3x**.

TECHNICAL SKILLS

- **AI/ML & Inference:** LLM inference/serving, prompt engineering, model evaluations, safety orchestration, RAG
- **Languages & Systems:** Python, C#, C++, KQL, OpenAI API, Kubernetes, Redis, Cosmos DB, Service Bus, .NET
- **Cloud & Infrastructure:** Azure (ARM, AKS, Service Bus, Cosmos DB), AWS (Bedrock, S3, Textract, ECR), ARM Templates, Azure DevOps, Grafana

SELECTED PROJECTS

- **AI Lab Assistant** (TurnTheBus nonprofit, 2025–Present): Architected and built end-to-end on **AWS** (Bedrock/Claude Haiku, S3, Textract, ECR); uses **RAG** over physics lab manuals to generate Socratic hints that guide high school students through virtual experiments based on live experiment state.

TECHNICAL WRITING

- *Runtime Optimizations for LLMs* (2026). Production notes on model offloading, KV cache management, and quantization for latency/cost-efficient LLM serving. pranshug.com/blog/runtime-optimizations-for-llms

EDUCATION

- **Georgia Institute of Technology** / *M.S. Computer Science (Machine Learning)* / Atlanta, GA May 2020
- **Indian Institute of Technology, Kanpur** / *B.Tech, C.S. and Eng.* / Kanpur, India May 2017